Algorithm 1 Construction of a set $\mathcal{A}$ with $H(\mathbf{c}_\mathcal{A}) < ML$ for a shift-XOR LRC

1: Set $\mathcal{A}_0 = \emptyset$ and $i = 1$
2: **while** $H(\mathbf{c}_{\mathcal{A}_{i-1}}) < ML$ **do**
3:     Pick $j_i \notin \mathcal{A}_{i-1}$ s.t. $|\Gamma(j_i) \setminus \mathcal{A}_{i-1}| \geq \delta - 1$
4:     **if** $H(\mathbf{c}_{\mathcal{A}_{i-1}}, \mathbf{c}_{\Gamma(j_i)}) < ML$ **then**
5:         set $\mathcal{A}_i = \mathcal{A}_{i-1} \cup \Gamma(j_i)$
6:     **else if** $H(\mathbf{c}_{\mathcal{A}_{i-1}}, \mathbf{c}_{\Gamma(j_i)}) \geq ML$ **then**
7:         $\mathcal{T} = \underset{\mathcal{T}' \subset \Gamma(j_i); H(\mathbf{c}_{\mathcal{A}_{i-1}}, \mathbf{c}_{\mathcal{T}'}) < ML}{\arg\max} |\mathcal{T}'|$
8:         **if** $\mathcal{T} = \emptyset$ **then**
9:             **break**
10:         **else**
11:             set $\mathcal{A}_i = \mathcal{A}_{i-1} \cup \mathcal{T}$
12:             **break**
13:         **end if**
14:     **end if**
15:     $i = i + 1$
16: **end while**
17: Output: $\mathcal{A} = \mathcal{A}_{i-1}$

Proof of Theorem 1. Let file $\mathbf{X}$ consist of M sequences, where each sequence contains L bits that are independent and identically distributed (i.i.d.) uniform random variables over the binary field. The (binary) entropy of this file is

$$H(\mathbf{X}) = ML. \tag{3}$$

For a general shift-XOR code that encodes the file $\mathbf{X}$ into n code blocks $\mathbf{C} = [\mathbf{c}_1, \mathbf{c}_2, \ldots, \mathbf{c}_n]$, where by the independence bound, each coded block has entropy

$$H(\mathbf{c}_i) \leq \alpha L + S. \tag{4}$$

By the defintion of (r, δ) locality, for any indices $j \in [n]$, there exists a set of indices $\Gamma(j) \subseteq [n]$, where $j \in \Gamma(j)$, $|\Gamma(j)| \leq r + \delta - 1$ and any code block indexed by this group can be reconstructed from any other r blocks in the same group. Then, we can deduce that $H(\mathbf{c}_{\Gamma(j)}) \leq r(\alpha L + S)$, where $\mathbf{c}_{\Gamma(j)} \triangleq \{\mathbf{c}_i : i \in \Gamma(j)\}$.

As the origianl file can be reconstructed from any k nodes, for any subset $\mathcal{A} \subseteq [n]$ with $|\mathcal{A}| \geq k$, we have $H(\mathbf{c}_\mathcal{A}) = ML$. Then we have

$$k \geq 1 + \max_{H(\mathbf{c}_\mathcal{A}) < ML} |\mathcal{A}|. \tag{5}$$

To establish a lower bound on k, we follow the approach in [1], [9], [21] to construct the largest possible subset $\mathcal{A} \subseteq [n]$ of indices that the corresponding blocks cannot reconstruct the original file. It is achieved by adapting an algorithm from [9], [21], as given in Algorithm 1. The algorithm starts with $\mathcal{A}_0 = \emptyset$, and iteratively expands $\mathcal{A}_{i-1}$ to $\mathcal{A}_i$ by adding at most $r + \delta - 1$ elements. For each iteration i, we define:

$$a_i = |\mathcal{A}_i| - |\mathcal{A}_{i-1}| \tag{6}$$

$$h_i = H(\mathbf{c}_{\mathcal{A}_i}) - H(\mathbf{c}_{\mathcal{A}_{i-1}}) \tag{7}$$

Let ℓ denote the number of iterations before termination, i.e., $\mathcal{A} = \mathcal{A}_\ell$. Then the final set size and entropy can be expressed as the sum of their increments:

$$|\mathcal{A}| = |\mathcal{A}_\ell| = \sum_{i=1}^{\ell} a_i \tag{8}$$

$$H(\mathbf{c}_\mathcal{A}) = H(\mathbf{c}_{\mathcal{A}_\ell}) = \sum_{i=1}^{\ell} h_i \tag{9}$$

In step 3, such a j_i always exists, since otherwise, by the (r, δ) locality, $\mathbf{c}_{\mathcal{A}_{i-1}}$ can recover all other coded blocks not indexed in $\mathcal{A}_{i-1}$ and hence $H(\mathbf{c}_{\mathcal{A}_{i-1}}) = ML$, which is a contradiction. Therefore, the algorithm can terminate at either step 9 or step 12, which are analyzed separately:

Case 1: Suppose that the algorithm terminates at step 9, i.e., after adding $\Gamma(j_\ell)$ to $\mathcal{A}_{\ell-1}$ in the ℓ-th iteration. For each iteration $i \leq \ell$, we analyze the entropy increment h_i as follows:

$$\begin{aligned}
h_i &= H(\mathbf{c}_{\mathcal{A}_i}) - H(\mathbf{c}_{\mathcal{A}_{i-1}}) \\
&= H(\mathbf{c}_{\mathcal{A}_{i-1} \cup (\mathcal{A}_i \setminus \mathcal{A}_{i-1})}) - H(\mathbf{c}_{\mathcal{A}_{i-1}}) \\
&= H(\mathbf{c}_{\mathcal{A}_{i-1}}) + H(\mathbf{c}_{\mathcal{A}_i \setminus \mathcal{A}_{i-1}} | \mathbf{c}_{\mathcal{A}_{i-1}}) - H(\mathbf{c}_{\mathcal{A}_{i-1}}) \\
&= H(\mathbf{c}_{\mathcal{A}_i \setminus \mathcal{A}_{i-1}} | \mathbf{c}_{\mathcal{A}_{i-1}}) \\
&\leq (a_i - \delta + 1)(\alpha L + S).
\end{aligned} \tag{10}$$

The last inequality holds as follow: in iteration i, $a_i = |\Gamma(j_i) \setminus \mathcal{A}_{i-1}| \geq \delta - 1$ new blocks are added to $\mathcal{A}_{i-1}$. Among the a_i newly added blocks in iteration i, at least $\delta - 1$ blocks are deterministic functions of the other blocks indexed in $\Gamma(j_i)$. Therefore, these $\delta - 1$ blocks do not contribute to the entropy increment, leaving at most $(a_i - \delta + 1)$ blocks that can contribute at most $(\alpha L + S)$ each to the entropy.

From (10), we can derive a lower bound on the size increment:

$$a_i \geq \frac{h_i}{\alpha L + S} + \delta - 1. \tag{11}$$

Substituting this into (8), we obtain:

$$\begin{aligned}
|\mathcal{A}| = |\mathcal{A}_\ell| &= \sum_{i=1}^{\ell} a_i \\
&\geq \sum_{i=1}^{\ell} \left(\frac{h_i}{\alpha L + S} + \delta - 1 \right) \\
&= \frac{1}{\alpha L + S} \sum_{i=1}^{\ell} h_i + (\delta - 1)\ell \tag{12} \\
&= \frac{1}{\alpha L + S} H(\mathbf{c}_\mathcal{A}) + (\delta - 1)\ell \tag{13}
\end{aligned}$$

We now give a lower bound on $H(\mathbf{c}_\mathcal{A})$ and ℓ. Since the algorithm is exiting, we know that no additional index $j \notin \mathcal{A}$ can be added to our current index set $\mathcal{A}$ without violating $H(\mathbf{c}_\mathcal{A}) < ML$. This means that $H(\mathbf{c}_\mathcal{A})$ must be sufficiently close to ML. Specifically, we argue that

$$H(\mathbf{c}_\mathcal{A}) \geq ML - (\alpha L + S). \tag{14}$$

This holds because if we assume otherwise, i.e., $H(\mathbf{c}_{\mathcal{A}}) \leq ML - (\alpha L + S) - \epsilon$ for any $\epsilon > 0$, then we could always find a new index $j \notin \mathcal{A}$ such that adding the corresponding coded block $\mathbf{c}_j$ would increase the total entropy by at most $(\alpha L + S)$, keeping the total entropy below ML. This contradicts the exit condition of the algorithm.

For the lower bound on ℓ, we analyze the entropy increment in each iteration. From the locality constraint, we know $a_i - \delta + 1 \leq r$, which together with (10) implies that $h_i \leq r(\alpha L + S)$. This means each iteration can contribute at most $r(\alpha L + S)$ to the total entropy increment. Furthermore, from (14) we know that the total entropy increment must reach at least $ML - (\alpha L + S)$. Therefore:

$$\ell \geq \left\lceil \frac{H(\mathbf{c}_{\mathcal{A}})}{r(\alpha L + S)} \right\rceil \geq \left\lceil \frac{M}{r(\alpha + \frac{S}{L})} \right\rceil - 1 \qquad (15)$$

Combining (13), (14), and (15) yields:

$$|\mathcal{A}_\ell| \geq \left\lceil \frac{M}{\alpha + \frac{S}{L}} \right\rceil - 1 + \left(\left\lceil \frac{M}{r(\alpha + \frac{S}{L})} \right\rceil - 1 \right)(\delta - 1). \quad (16)$$

Case 2: Here, the algorithm terminates at step 12 in the ℓ-th iteration because:

$$H(\mathbf{c}_{\mathcal{A}_{\ell-1} \cup \Gamma(j_\ell)}) \geq ML. \qquad (17)$$

Since each iteration increases the entropy by at most $r(\alpha L + S)$, we have:

$$\ell \geq \left\lceil \frac{ML}{r(\alpha L + S)} \right\rceil. \qquad (18)$$

The size increments are bounded as in *Case 1*:

$$a_i \geq \frac{h_i}{\alpha L + S} + \delta - 1, \quad \text{for } i \leq \ell - 1. \qquad (19)$$

Note that in the last iteration, we only add a subset $\mathcal{T} \subset \Gamma(j_\ell)$ rather than the set, thus we cannot guarantee $|\mathcal{T}| \geq \delta - 1$ and the existence of $\delta - 1$ deterministic blocks. For the last iteration,

$$a_\ell \geq \frac{h_\ell}{\alpha L + S}. \qquad (20)$$

Combining (8), (14), (18), (19), and (20), we obtain:

$$\begin{aligned} |\mathcal{A}_\ell| &= \sum_{i=1}^{\ell} a_i \\ &\geq \sum_{i=1}^{\ell-1} \left(\frac{h_i}{\alpha L + S} + \delta - 1 \right) + \frac{h_\ell}{\alpha L + S} \\ &= \frac{1}{\alpha L + S} H(\mathbf{c}_{\mathcal{A}}) + (\ell - 1)(\delta - 1) \\ &\geq \left\lceil \frac{M}{\alpha + \frac{S}{L}} \right\rceil - 1 + \left(\left\lceil \frac{M}{r(\alpha + \frac{S}{L})} \right\rceil - 1 \right)(\delta - 1). \end{aligned} \quad (21)$$

Finally, combining (5), (16), and (21), we establish the lower bound:

$$k \geq \left\lceil \frac{M}{\alpha + \frac{S}{L}} \right\rceil + \left(\left\lceil \frac{M}{r(\alpha + \frac{S}{L})} \right\rceil - 1 \right)(\delta - 1). \qquad \square$$

Proof of Theorem 2. Each sequence with subscript j in the fail node of repair group $i, i \in [N]$ is contained in a $(\delta - 1, r)$ shift-XOR system

$$[p_j^1, p_j^2, \ldots, p_j^{\delta-1}]^T = \mathbf{\Omega}[y_j^1, \ldots, y_j^r]^T,$$

where $j \in [(i-1)(r+\delta-1), i(r+\delta-1)], i \in [N]$. We have $\sum_{(i-1)(r+\delta-1)}^{i(r+\delta-1)} \sigma_j = r(\delta - 1)$, which is independent of helper node selection. The recovery of sequences in $\mathbf{Y}$ requires $(r-1)r(\delta-1)L$ XOR operations. Generating local parity sequences requires an additional $(r-1)(\delta-1)(L + O(r\delta))$ XOR operations, totaling $(r^2-1)(\delta-1)(L + O(r\delta))$ XOR operations. $\square$

Proof of Theorem 3. As $B = \left\lceil \frac{rk}{r_\delta} \right\rceil + \left(\left\lceil \frac{k}{r_\delta} \right\rceil - 1 \right)(\delta - 1)$, and $k = r_\delta(q-1) + t$ where $1 \leq t < r_\delta$, we have:

$$\begin{aligned} B - k &= r(q-1) + \left\lceil \frac{rt}{r_\delta} \right\rceil + (q-1)(\delta-1) - r_\delta(q-1) - t. \\ &= \left\lceil \frac{rt}{r_\delta} \right\rceil - t. \end{aligned}$$

Since $\frac{rt}{r_\delta} = \frac{rt}{r+\delta-1} < t$, we have $B - k \leq 0$.

On the other hand, since $\left\lceil \frac{rt}{r_\delta} \right\rceil \geq \frac{rt}{r_\delta}$ and $1 \leq t < r_\delta$:

$$B - k \geq -t \left(\frac{\delta - 1}{r_\delta} \right) \geq -(r_\delta - 1)\left(\frac{\delta - 1}{r_\delta} \right) > -(\delta - 1).$$

Therefore, $-(\delta - 1) < B - k \leq 0$. Since both B and k are integers:

$$|B - k| \leq \delta - 2.$$

For $B = k$ to hold, we require:

$$\left\lceil \frac{rt}{r_\delta} \right\rceil = t \Leftrightarrow \frac{rt}{r_\delta} > t - 1 \Leftrightarrow r > (\delta-1)(t-1)$$

When $\delta = 2$, $k = B$ always holds. $\square$

Proof of Theorem 4. When $b = 0$, for each codeword, accessing any r nodes from each of the q parity groups allows recovery of all rg^F global parity sequences and computation of new global parity sequences. This requires accessing λqr nodes in the initial codewords and g^F nodes in the final codeword for data writing.

When $b \neq 0$ and $g^F \leq b$, all required global parity sequences are in the mixed group. By accessing any r nodes in the mixed group, we need λr node accesses in the initial codewords. Additionally, regrouping the systematic sequences in the mixed group of each initial codeword requires accessing $\lambda(r_\delta - d)$ nodes in the final codeword.

When $b \neq 0$ and $g^F > b$: If $h = 0$, the required global parity sequences span over $(q+1)$ repair groups, requiring $\lambda(q+1)r$ node accesses in the initial codewords, plus $\lambda(r_\delta - b)$ nodes in the final codeword for regrouping.

If $h > 0$, the required global parity sequences span over $(q+2)$ repair groups, necessitating $\lambda(q+2)r$ node accesses in the initial codewords, plus $\lambda(r_\delta - b)$ nodes in the final codeword for regrouping. Due to the cyclic shift pattern, collecting global

parity sequences with the same subscript requires accessing r nodes within a repair group, making this strategy consistently more efficient than direct access.

In all cases, we need additional g^F nodes access for writing the new global parity sequences. □

Proof of Theorem 5. From [17], for an (n, k) Two-Tone shift-XOR system, the storage overhead is at least

$$\begin{cases} \frac{(n^2-1)(k-1)}{4}, & \text{if } n \text{ is odd} \\ \frac{n^2(k-1)}{4}, & \text{if } n \text{ is even} \end{cases}$$

with equality if and only if using a reflected Vandermonde matrix.

For the global parity generation, matrix $\mathbf{\Phi}$ can be viewed as r separate $(n-k, k)$ shift-XOR systems. Each system has a storage overhead of at least:

$$S(\mathbf{\Phi}) = (k-1)\eta_1$$

where $\eta_1 = \frac{(n-k)^2 - \alpha_1}{4}$ and $\alpha_1 = 1$ if $n-k$ is odd (0 if even).

Similarly, the local parity generation through matrix $\mathbf{\Omega}$ can be viewed as n separate $(\delta-1, r)$ shift-XOR systems, where

$$S(\mathbf{\Omega}) = (r-1)\eta_2$$

where $\eta_2 = \frac{(\delta-1)^2 - \alpha_2}{4}$ and $\alpha_2 = 1$ if $\delta - 1$ is odd (0 if even).

Since the local parity generation operates on shifted sequences from $\mathbf{Y}$, it introduces an extra storage overhead of $(\delta - 1)S(\mathbf{\Phi})$.

Therefore, the total storage overhead is $S' = (r + \delta - 1)S(\mathbf{\Phi}) + nS(\mathbf{\Omega})$, which equals $\hat{S}$ if and only if both $\mathbf{\Phi}$ and $\mathbf{\Omega}$ are Reflected Vandermonde matrices. □